

Desarrollo de un programa de computador capaz de jugar Starcraft®: Brood War™ usando técnicas de aprendizaje por refuerzo

Diego A. Ballesteros Villamizar

Universidad de los Andes

October 27, 2011

Outline

- 1 **RL Library**
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 **Scenario with incomplete game information**
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 **Test Case #3**
 - Introduction
 - Implementation
- 4 **To Do List**

Outline

- 1 **RL Library**
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

Motivation

- During the development of this project several scenarios were explored, each with a different model
- But the algorithms remained the same, then some code was copy-pasted among projects
- So it is convenient to create a library that can be adapted to different problems and implements the reinforcement learning algorithms

Similar Work

- A framework for reinforcement learning is not something new
- RL-Glue is a software that provides a language-independent interface to connect different RL software blocks ¹
- RL-Toolbox is a C++ library that implements several algorithms and a extensive framework that includes the agent, controller and environment. Has a special focus on optimal control of continuous tasks ²

¹Tanner, B., & White, A. (2009). RL-Glue : Language-independent software for reinforcement-learning experiments. *Journal of Machine Learning Research*, 10, 2133–2136

²Neumann, G. (2005). *The Reinforcement Learning Toolbox, Reinforcement Learning for Optimal Control Tasks*. Master's thesis, Technische Universität Graz, Austria

So Why Implement It?

- RL-Toolbox is a great option and it is in C++ but not updated since 2006
- To gain better insight on reinforcement learning techniques I consider it is better to implement them
- Also due to time constraints it is a safer path that understanding and extending another framework
- Setup a start point that can be extended by other people interested in reinforcement learning, a simpler yet useful RL framework for education

Outline

- 1 **RL Library**
 - Introduction
 - **Design**
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

A View of the Reinforcement Learning Framework

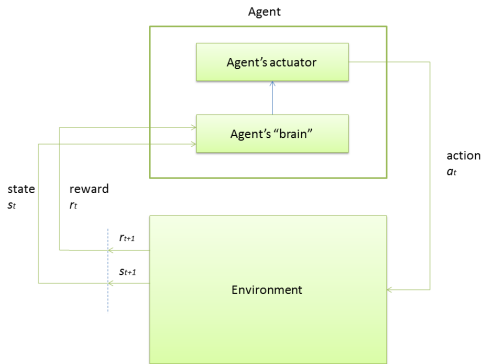


Figure: Depiction of the proposed RL Framework

Basic Reinforcement Learning Process

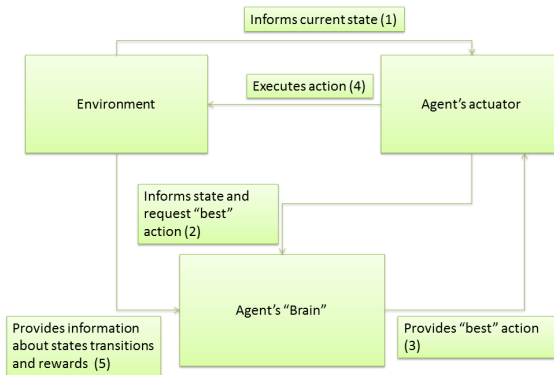


Figure: Illustration of the basic RL process, the numbers indicate the actions' order

Library Design

Considerations

- The library will represent the agent's "brain"
- Must process the environment information, implements a learning algorithm and indicates the actions to take in every state
- Should define a general interface so it can be integrated with multiple problems, represented by an environment (states) and an actuator (actions) as well as state transitions (environment's dynamics) and a reward model

Library Design

Scope

- For now only discrete reinforcement learning problems are considered
- Only model free TD(0) algorithms will be implemented (Q-Learning and SARSA)
- Most of the work consists in the implementation of a general framework for discrete problems



Figure: Trying to avoid the scope creep

Library Design

Problem Representation

- Problems can be represented as a collection of states and actions in every state (state-action pairs)
- States can be represented as sets of multi-valued features otherwise states must have an unique label
- Actions can be represented with an unique label
- Use of states represented by features can help fast search of state-action pairs, by indexing them using the features' values

Library Design

Algorithms' Data Representation

- Once a problem has been represented in terms of its state-action pairs then learning information must be stored too
- For Q-Learning and SARSA the only needed information are Q-values
- A lexicographical order of state-action pairs is performed and with this they are indexed for faster search
- Q-Values can be stored in an array indexed the same way as state-action pairs
- Features can be used to build indexes and have instant access to the information

Library Design

Information Storage

- Problems can be stored in text files and loaded to the library during execution
- Q-Values can be stored in text files, and loaded posteriorly to continue the learning
- Resulting policies can also be generated and stored in a text file
- These files' format will be shown shortly

Outline

- 1 **RL Library**
 - Introduction
 - Design
 - **Implementation**
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

Class Diagram

Problem Representation

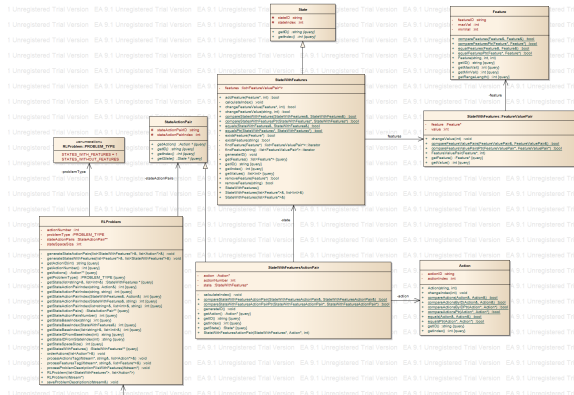


Figure: Classes used to describe a RL problem

Class Diagram

Algorithms' implementation

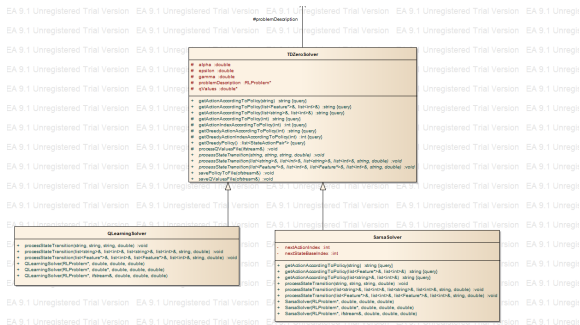


Figure: Classes that implement the reinforcement learning algorithms

Problem Description File

```
#Lines started with '#' are comments lines and are ignored
#The first line to be processed must be the type line
#Therefore every line must be either a comment line or
#a line with proper formatting
#For the formats following in the comments, a SSSS represents a string
# and a ### represents a positive integer
#The type line has the following format
#Type=###
#1 represents a problem defined with states, defined by features, and actions
#2 represents a problem with states, without features, and actions
Type=1
#After type comes the problem description
#The next non-comment line must be a tag
#Possible tags are:
#[Features] only if type was set to 1
#[States] only if type was set to 2
#[Actions] for any type
[Features]
#Each feature is described in a single line, the description must be written
#in the following format
#ID=SSSS;Min=###;Max=###
ID=X;Min=0;Max=4
ID=Y;Min=0;Max=4
#After all items are described the next non-comment line must be one
#of the remaining tags depending on the type of problem description selected
#Tags must not be repeated
[Actions]
#Following an actions tag comes the definition of all available actions
#Each definition must be in the following format
#Action=SSSS
Action=UP
Action=DOWN
Action=LEFT
Action=RIGHT
```

```
#This file describes the states and actions for the windyGridWorld
Type=1
[Features]
ID=X-Coordinate;Min=0;Max=9
ID=Y-Coordinate;Min=0;Max=6
[Actions]
Action=UP
Action=DOWN
Action=LEFT
Action=RIGHT
```

Figure: An example of a problem description file

Figure: General file's format

Q-Values File

```
#A comment line is preceded by a '#' symbol
#In every automatically generated file the first line
#will be the following comment line:
#"This file has been automatically generated"
#Next there will be a line indicating the type o
#f information contained in the file, in this case:
#"This file contains learned Q-values"
#Next comes the type of problem, that is either 1 or 2
Type=#
#The next line is a tag indicating the start of
#the Q-Values information
[Q-Values]
#Then each line will contain a state ID
#separated by a ':' from the action and
#then another ';' before the associated Q-Value
STATE 1 ; MOVE ; Q-Value= 0.32
STATE 1 ; DANCE ; Q-Value= 0.40
```

Figure: General file's format

```
#This file has been automatically generated
#This file contains learned Q-values
Type=1
[Q-Values]
X-Coordinate= 0 ; Y-Coordinate= 0 ; DOWN ; Q-Value= -6.26433
X-Coordinate= 0 ; Y-Coordinate= 0 ; LEFT ; Q-Value= -6.28287
X-Coordinate= 0 ; Y-Coordinate= 0 ; RIGHT ; Q-Value= -6.19639
X-Coordinate= 0 ; Y-Coordinate= 0 ; UP ; Q-Value= -6.20517
X-Coordinate= 0 ; Y-Coordinate= 1 ; DOWN ; Q-Value= -6.57774
X-Coordinate= 0 ; Y-Coordinate= 1 ; LEFT ; Q-Value= -6.55397
X-Coordinate= 0 ; Y-Coordinate= 1 ; RIGHT ; Q-Value= -6.54256
X-Coordinate= 0 ; Y-Coordinate= 1 ; UP ; Q-Value= -6.57537
X-Coordinate= 0 ; Y-Coordinate= 2 ; DOWN ; Q-Value= -7.13986
X-Coordinate= 0 ; Y-Coordinate= 2 ; LEFT ; Q-Value= -7.266
X-Coordinate= 0 ; Y-Coordinate= 2 ; RIGHT ; Q-Value= -7.13619
X-Coordinate= 0 ; Y-Coordinate= 2 ; UP ; Q-Value= -7.11861
X-Coordinate= 0 ; Y-Coordinate= 3 ; DOWN ; Q-Value= -7.82849
X-Coordinate= 0 ; Y-Coordinate= 3 ; LEFT ; Q-Value= -7.93717
X-Coordinate= 0 ; Y-Coordinate= 3 ; RIGHT ; Q-Value= -7.84577
X-Coordinate= 0 ; Y-Coordinate= 3 ; UP ; Q-Value= -7.82417
X-Coordinate= 0 ; Y-Coordinate= 4 ; DOWN ; Q-Value= -7.77527
X-Coordinate= 0 ; Y-Coordinate= 4 ; LEFT ; Q-Value= -7.85447
X-Coordinate= 0 ; Y-Coordinate= 4 ; RIGHT ; Q-Value= -7.78909
X-Coordinate= 0 ; Y-Coordinate= 4 ; UP ; Q-Value= -7.82938
```

Figure: An example of a Q-Values file

Policy File

```
#A comment line is preceded by a '#' symbol
#In every automatically generated file the first line
#will be the following comment line:
#"This file has been automatically generated"
#The next line is a tag indicating the start of
#the policy information
[Policy]
#Then each line will contain a state ID
#separated by a ':' from the best action found after
#learning
STATE 1 ; DANCE
STATE 2 ; MOVE
```

Figure: General file's format

```
#This file has been automatically generated
[Policy]
X-Coordinate= 0 ; Y-Coordinate= 0 ; RIGHT
X-Coordinate= 0 ; Y-Coordinate= 1 ; RIGHT
X-Coordinate= 0 ; Y-Coordinate= 2 ; UP
X-Coordinate= 0 ; Y-Coordinate= 3 ; UP
X-Coordinate= 0 ; Y-Coordinate= 4 ; DOWN
X-Coordinate= 0 ; Y-Coordinate= 5 ; UP
X-Coordinate= 0 ; Y-Coordinate= 6 ; LEFT
X-Coordinate= 1 ; Y-Coordinate= 0 ; UP
X-Coordinate= 1 ; Y-Coordinate= 1 ; LEFT
X-Coordinate= 1 ; Y-Coordinate= 2 ; UP
X-Coordinate= 1 ; Y-Coordinate= 3 ; DOWN
X-Coordinate= 1 ; Y-Coordinate= 4 ; DOWN
X-Coordinate= 1 ; Y-Coordinate= 5 ; DOWN
X-Coordinate= 1 ; Y-Coordinate= 6 ; RIGHT
X-Coordinate= 2 ; Y-Coordinate= 0 ; DOWN
X-Coordinate= 2 ; Y-Coordinate= 1 ; RIGHT
X-Coordinate= 2 ; Y-Coordinate= 2 ; UP
X-Coordinate= 2 ; Y-Coordinate= 3 ; DOWN
X-Coordinate= 2 ; Y-Coordinate= 4 ; RIGHT
X-Coordinate= 2 ; Y-Coordinate= 5 ; LEFT
X-Coordinate= 2 ; Y-Coordinate= 6 ; LEFT
X-Coordinate= 3 ; Y-Coordinate= 0 ; RIGHT
X-Coordinate= 3 ; Y-Coordinate= 1 ; RIGHT
```

Figure: An example of a policy file

Outline

- 1 **RL Library**
 - Introduction
 - Design
 - Implementation
 - **Tests**
- 2 **Scenario with incomplete game information**
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 **Test Case #3**
 - Introduction
 - Implementation
- 4 **To Do List**

Testing on a Windy Gridworld

Windy Gridworld Problem




Figure: Image of the windy GridWorld problem ³

³Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning : an introduction*. A Bradford Book. The MIT Press

Testing on a Windy Gridworld

Comparison with Sutton & Barto's Book




Figure: Obtained results using the implemented library




Figure: Reinforcement Learning book's results

Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning : an introduction*. A Bradford Book. The MIT Press

Testing on a Windy Gridworld

Comparison and effect of learning rate




Figure: Obtained results using the implemented library




Figure: Results obtained by a student at Stellenbosch University

Schaaf, H. (2011). The basics of reinforcement learning.

URL <http://ml.sun.ac.za/2010/06/18/>

the-basics-of-reinforcement-learning/

Testing on the Cliff-Walking Problem

Cliff Walking Problem




Figure: Depiction of the Cliff-Walking problem ⁴

⁴Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning : an introduction*. A Bradford Book. The MIT Press

Testing on the Cliff-Walking Problem

Comparing Q-Learning and SARSA




Figure: Obtained results using the implemented library



Figure: Reinforcement Learning book's results

Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning : an introduction*. A Bradford Book. The MIT Press

Outline

- 1 RL Library
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

Map Details

- Open map, no obstacles
- 128x128 Tiles (4096x4096 pixels)



Figure: Map editor view


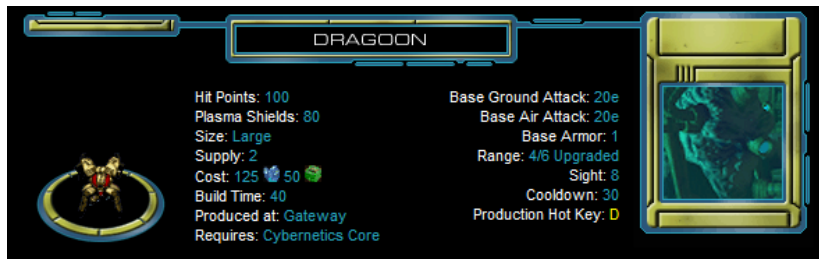


Figure: Minimap view

Unit Details

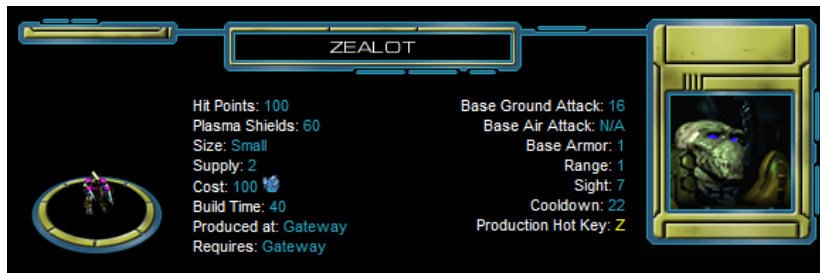


The image shows the Dragoon unit information screen from the game StarCraft II. At the top, the unit name "DRAGOON" is displayed in a blue-bordered box. Below this, the unit's stats are listed in two columns. On the left, there is a small icon of the Dragoon unit. On the right, there is a small image of the unit's production building, the Gateway. The stats are as follows:

Stat	Value
Hit Points	100
Plasma Shields	80
Size	Large
Supply	2
Cost	125 50
Build Time	40
Produced at	Gateway
Requires	Cybernetics Core
Base Ground Attack	20e
Base Air Attack	20e
Base Armor	1
Range	4/6 Upgraded
Sight	8
Cooldown	30
Production Hot Key	D

Figure: Dragoon unit information

Unit Details



The image shows the unit information screen for a Zealot in StarCraft. At the top, the unit name "ZEALOT" is displayed in a blue-bordered box. Below this, the unit's stats are listed in two columns. On the left, there is a small icon of a Zealot standing on a circular platform. On the right, there is a larger image of a Zealot's head in a yellow frame. The stats are as follows:

Stat	Value
Hit Points	100
Plasma Shields	60
Size	Small
Supply	2
Cost	100
Build Time	40
Produced at	Gateway
Requires	Gateway
Base Ground Attack	16
Base Air Attack	N/A
Base Armor	1
Range	1
Sight	7
Cooldown	22
Production Hot Key	Z

Figure: Zealot unit information

Unit Details

- Bot: 1 Dragoon
- Enemy: 1 Zealot
- A Zealot can do 0.72 Damage per second against a Dragoon
- A Dragoon can do 0.33 Damage per second against a Zealot
- 12 Hits from a Zealot kill a Dragoon, that would take 264 frames (ca. 11 seconds)
- 16 Hits from a Dragoon kill a Zealot, that would take 480 frames (ca. 20 seconds)
- The zealot has the upper hand

Unit Details

- Bot: 1 Dragoon
- Enemy: 1 Zealot
- A Zealot can do 0.72 Damage per second against a Dragoon
- A Dragoon can do 0.33 Damage per second against a Zealot
- 12 Hits from a Zealot kill a Dragoon, that would take 264 frames (ca. 11 seconds)
- 16 Hits from a Dragoon kill a Zealot, that would take 480 frames (ca. 20 seconds)
- **Maybe Not?**

Outline

- 1 RL Library
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - **Introduction**
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

What changes?

- So far the reinforcement learning agent has cheated!
- A complete map information flag was enabled and the agent knew every detail of the game
- So how does the agent perform if some information is hidden?

Motivation

- Real world problems do not necessarily have all the information available at all times
- In Starcraft most information about the enemy is usually hidden to the player
- Fog of War hides units' position and status
- Implementing working algorithms in this environment is important

Outline

- 1 RL Library
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - **Implementation and Results**
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

How to deal with the new condition?

- Partial Observable Markov Decision Processes describe this kind of problem
- There literature describes algorithms but these are rather complicated
- For this game it may be possible to predict the states based on memory and knowledge of the game

How to deal with the new condition?

- In this particular scenario it may not be relevant because when the enemy is hidden it doesn't pose a threat
- The enemy always look for the Dragoon so he can just sit and wait until the environment is observable again
- When the enemy is out of sight, the Dragoon moves randomly around its position

Learning results




Figure: Winning probability smoothed using a moving average filter

Analysis

- Results are similar to those of the complete information case
- As the enemy always pursues the agent, the enemy is visible most of the time
- It could be interesting to explore the capabilities of using specific knowledge of the game to deal with partial information in a reinforcement learning problem

Outline

- 1 RL Library
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - Implementation
- 4 To Do List

Motivation

- The hardest of all experiments so far, multiple units on both sides and a completely balanced match
- Developing a good behavior that can help the player attack a position is a step in the right direction to the development of AIs using reinforcement learning
- There are countless possibilities and it is interesting to see what reinforcement learning can produce
- This experiment is also a good scenario to implement a conventional FSM-based AI for comparison

Test case's map




Figure: Agent's Units

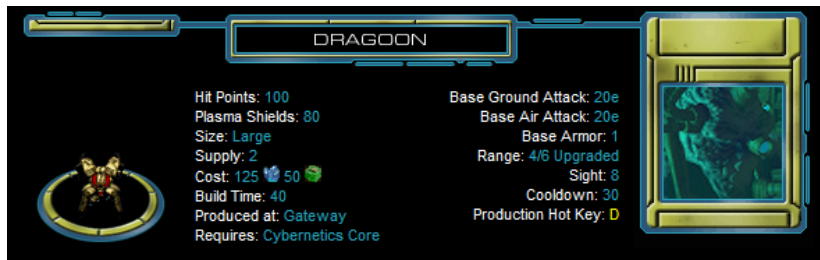


Figure: Enemy's Units

Units' Details

- 6 Dragoons for each player
- The agent's dragoons will be in one corner of the map
- The enemy's dragoons will remain in the opposite corner and will be defending the zone
- The game is balanced and micromanagement decides the outcome

Unit's Details



The image shows the Dragoon unit information screen from StarCraft II. At the top, the unit name "DRAGOON" is displayed in a blue-bordered box. Below this, the unit's stats are listed in two columns. On the left, there is a small icon of the Dragoon unit. On the right, there is a small image of the unit's production building, the Gateway. The stats are as follows:

Stat	Value
Hit Points	100
Plasma Shields	80
Size	Large
Supply	2
Cost	125 50
Build Time	40
Produced at	Gateway
Requires	Cybernetics Core
Base Ground Attack	20e
Base Air Attack	20e
Base Armor	1
Range	4/6 Upgraded
Sight	8
Cooldown	30
Production Hot Key	D

Figure: Dragoon unit information

Objective

- Control the group of dragoons and move them to the objective and destroy any enemies encountered

Outline

- 1 RL Library
 - Introduction
 - Design
 - Implementation
 - Tests
- 2 Scenario with incomplete game information
 - A Quick Recap
 - Introduction
 - Implementation and Results
- 3 Test Case #3
 - Introduction
 - **Implementation**
- 4 To Do List

Actions

Reinforcement Learning Setup

- Move to Target
- Retreat from Target
- Get in some formation
- Attack weakest unit in range
- Attack weakest unit in group range
- Attack weakest unit in sight
- Attack weakest unit in group sight
- Retreat units under attack
- Retreat weakest units under attack

Formations

Actions

- Form a straight line: To cover more area and attack multiple enemies at once
- Form an arc: To focus fire in tightly grouped enemies
- Group: To move near a point before moving
- Disperse: Break formation to explore more territory

States

- Formation Status: Forming, In Formation (Which one?)
- Units Alive: In, discretized, %
- Enemies Killed: In, discretized, %
- Distance to Target
- Enemies in Group Range: In, discretized, %

Schedule

- Get results from third test case, November 13
- Prepare final presentation and document, Week of November 14
- Hand final document to supervisor, November 21
- Hand final document to reviewers, November 30
- Presentation, December 7 or 9